



Barndoor.ai

The 2026 Enterprise Guide For AI & MCP Security



2026 Strategic Guide For AI & MCP Security

Table Of Contents

01

2026 is All About
Making AI Work

02

The Model Context
Protocol (MCP)
Overview

03

How MCP Transforms
Enterprise AI

04

MCP Fundamentally
Changes the Enterprise
Risk Profile

05

Recommendations
For AI Agents & MCP
Security

06

Barndoor Provides
Secure Access For AI

07

Key Business Benefits
With Barndoor

08

Ready for 2026?



2026 Is All About Making AI Work

2025 was the year enterprises experimented with AI. Teams ran pilots, tested different AI clients, compared LLM models, spun up their own Model Context Protocol (MCP) servers, and connected to a growing ecosystem of public MCP servers. It was a year of exploration—figuring out what AI could do, how it behaved, and where it might fit.

2026 is the year AI must deliver. Enterprises now expect measurable productivity gains and real cost efficiency across every department. Finance, sales, engineering, operations—each needs AI that works reliably, securely, and at scale within existing business systems.

The question is no longer “*Can AI help?*”

It’s “*How do we make AI safe, compliant, and operationally efficient?*”



2025

Experimentation Phase

- Pilot projects
- Multiple AI clients
- LLM model switching
- DIY MCP servers
- Public MCP connections
- Prompt testing
- Low-stakes automation



2026

Operationalization Phase

- Enterprise-wide rollouts
- Standardized AI clients
- Approved LLM models
- Trusted MCP registry
- Governed data access
- Secure tool calls
- Multi-system use case

The Model Context Protocol (MCP) Overview

The Model Context Protocol (MCP) has become the de facto standard for connecting AI agents to business-critical data. Much like APIs transformed software integration in the early 2000s, MCP introduces a universal interface to contextualize enterprise data, whether it's a CRM, ticketing system, codebase, or document repository. MCP allows AI agents to plug into these business systems, interpret the user request, analyze available data, and then either take action or return relevant information.

MCP Operationalizes AI Across Business Systems

The result? Higher productivity, faster onboarding of AI agents, and scalable automation that can drive measurable impact across the organization.

At the core of every MCP server is a set of defined tool calls—standardized, structured requests that AI agents make to MCP servers, such as:



Querying
records



Updating
fields



Triggering
workflows



Creating tickets



Deleting assets

How MCP Transforms Enterprise AI

MCP isn't just about accessing information—it's about making enterprise data usable and actionable for AI.

With MCP, Confluence pages, Jira tickets, Salesforce records, Google Drive docs, Slack conversations, Microsoft emails, and Zoom call transcripts are no longer isolated data silos. Through tool calls, AI agents can read from these business systems, synthesize information across them, and take action.

MCP Enables Intelligent Actions Within Each Platform



Salesforce

Query pipeline, update records, flag at-risk opportunities for follow-up.



Google Drive

Retrieve recent docs, analyze contents, and share securely with teams.



Slack

Monitor threads, trigger alerts, and generate instant status updates.



Snowflake

Run SQL queries, compile BI dashboards, and refresh data pipelines.



Plaid

Validate account balances, assess transaction patterns, automate financial reporting.



GitHub

Review pull requests, detect code anomalies, trigger CI/CD workflows.

However, the true power of MCP is its ability to enable complex, multi-system use cases that go beyond simple querying or monitoring. By connecting disparate business systems, MCP transforms disconnected business operations into coordinated, outcome-driven workflows, improving productivity.

MCP Drives Business Impact With Multi-System Use Case



AI checks the calendar

Sees the **weekly sales update** is due in the next 24 hours

Pulls opportunity list from Salesforce

Filters for deals lacking recent engagement



Checks Gmail to find opportunities with no customer response

Draft follow-up email for each



Sends Slack message to AE to

review and approve email content



Creates a slide summarizing the status of each opportunity

MCP Fundamentally Changes The Enterprise Risk Profile

What makes MCP so powerful is also what makes it dangerous: anyone can create an MCP server. While many developers connect to official, vendor-supported servers—like those offered by Salesforce, Notion, Asana, GitHub, and Box—they may also access third-party MCP servers spun up by individuals, research projects, or other sources that are not vetted or reliably hosted.

A connection to an untrusted server can cascade into unauthorized actions, data exfiltration, or compliance violations. And when AI agents orchestrate workflows across multiple business systems simultaneously, the complexity and risk increase.

Risks

Shifting MCP Specifications

Frequent updates and shifting MCP standards disrupt stability. Teams spend cycles fixing compatibility issues, which increases maintenance overhead and delays deployments.

Risks

Untrusted Servers

Many MCP servers are not production-ready, lack proper documentation, or are built by unreliable sources that may run unsafe code and expose sensitive data.

Risks

Non-standard Deployments

MCP servers are deployed across cloud, on-prem, containers, and VMs—often with inconsistent configurations. This makes it difficult for security teams to monitor access, enforce policies, and maintain a consistent security posture across environments.

Risks

Visibility Gaps

Organizations lack insight into who is behind each AI agent, which MCP servers they connect to, and what business systems or data they attempt to access. Without this visibility, governance and oversight become nearly impossible.

Risks

Shadow AI

Dev teams deploy AI agents that connect to MCP servers without security approval or oversight, resulting in unsanctioned MCP traffic and unmanaged exposure points.

Risks

No Security Standardization

The lack of standardized security and access controls across different MCP servers makes them vulnerable to unsafe tool calls, prompt injections, and compliance risks.

Risks

Inconsistent AuthN and AuthZ

Authentication and authorization policies are inconsistently applied across AI agents and MCP servers, leaving gaps in enforcement. This fragmented approach makes it challenging to manage permissions and ensure least-privilege access.

Risks

Dependency on Legacy Tools

Traditional Identity and Access Management (IAM) solutions or API gateways were built for humans and APIs—not AI-to-data interactions. They extend user-level access to AI agents, which can lead to unauthorized actions and unintended data changes or deletions.

Recommendations For AI & MCP Agent Security

Most organizations have spent years building controls around these applications—user permissions, access policies, approval workflows across Salesforce, Google Drive, GitHub, and more. But when an AI agent executes multi-system workflows using tool calls, traditional application-level controls provided by IAM and API gateways aren't enough. You need visibility and AI-specific access controls to determine which tool calls are permitted, which business systems can interact, and what actions AI apps and agents are authorized to perform.

Recommendations

✓ — Prevent Shadow AI With A Central MCP Registry

As teams deploy in-house servers and connect to third-party providers, your MCP landscape will fragment, creating security and compliance blind spots. Establish a single, enterprise-wide registry for all approved MCP servers and enforce policies that block connections to any unvetted MCP servers.

✓ — Enforce Version Control For MCP Servers

Standardize and approve specific MCP server versions across environments. Block outdated, unstable, or untested versions that may introduce unexpected behaviors or compatibility risks.

✓ — Limit AI Agent Access To Trusted MCP

Restrict AI agents to connect to official, vendor-supported MCP servers or internally approved MCP servers, blocking access to community-hosted or unverified servers that may lack proper security, reliability, or documentation.

✓ — Create Distinct Identities For AI Apps & Agents

Provide AI agents with distinct identities, separate from human users, but have an identity binding between them to maintain traceability and avoid granting AI agents full user-level permissions by default.

✓ — Integrate With Enterprise Identity Providers (IdPs)

Ensure AI agent authentication and authorization are integrated with existing enterprise IdPs. Leverage centralized identity and SSO frameworks to enforce consistent policies, simplify human user-AI agent mapping, and reduce the risk of unmanaged or orphaned agent access.

✓ Implement Granular Access Controls

Apply fine-grained role-based access controls (RBAC) to each AI agent and MCP server. Restrict tool calls based on user, role, use case, business function, and risk level, and ensure AI agents operate under least privilege access at all times.

✓ Get Unified Visibility Into AI Activities

Track every AI agent activity—see who's behind each agent, what MCP servers they connect to, and what business systems and data they access. Enable centralized visibility across all AI agents and MCP servers to eliminate shadow AI.

✓ Block Unsafe Tool Calls

Allow safe queries and standard operations while restricting destructive actions, such as mass deletes or bulk updates, unless explicitly authorized via approval workflows—block access to sensitive data based on user role, business function, or regulatory requirements.

✓ Monitor And Log All AI Interactions

Track and record every AI agent interaction with MCP servers, capturing details of all requests, including tool calls, successful operations, failed attempts, and security events for compliance. Aggregate logs across all AI-related activities to enable real-time anomaly detection based on volume, frequency, or behavioral deviations.

✓ Establish An AI-Incident Response Plan


Enable teams to quickly disable or isolate AI agents and MCP servers in response to security incidents, suspicious behavior, or policy violations. This ensures fast containment without disrupting other AI workflows, giving teams time to investigate and remediate securely.

**Guardrails Needed,
Urgently** →

MCP is not just a protocol. It's a new frontier of productivity *and* risk. But without consistent security, visibility, and control, AI projects will stall before reaching scale.

Barndoor Provides Secure Access For AI

Barndoor is a data and access management platform purpose-built to secure how AI interacts with your business data and applications. It mitigates the most critical risks of unmanaged MCP deployments through centralized governance, consistent policy enforcement, and real-time oversight. With Barndoor, you can define fine-grained access control policies, enforce them automatically, and gain end-to-end visibility into every AI request—down to the user, role, AI client, and specific tool call.



How Barndoor Eliminates Key AI & MCP Security Risks

Controlled AI Access

With Barndoor, employees connect only to vetted, pre-approved MCP servers—not to random public ones. Administrators pre-register the AI agent, the MCP server, and approve permitted tool calls, reducing exposure to unreliable or unvetted sources and unexpected actions.

Eliminated Shadow AI

By enforcing centralized visibility and approval workflows, Barndoor eliminates shadow AI development. Every AI agent—its purpose, ownership, and access scope—is tracked and governed, preventing unsanctioned deployments that bypass corporate security policies.

Standardized Security

Barndoor applies consistent security policies across all MCP servers and AI agents. This eliminates the fragmented posture that arises when teams deploy MCP servers with inconsistent authentication, authorization, and access configurations.

Granular Access Controls

Enforce fine-grained permissions down to the user, role, and tool call—ensuring AI agents only execute authorized operations across your MCP servers and connected business systems.

Real-Time Evaluation

Validate every request in real time by inspecting AI agent identity, tool call type, and business intent—before it's allowed to run.

Unified Monitoring & Visibility

Every MCP connection is tracked, logged, analyzed, and flagged, providing real-time visibility into what's happening. Provides insights on which MCP servers are most used via which AI agent, what sort of requests are being made, and what business systems or data are accessed.

Key Business Benefits With Barndoor



Faster Time To Value

Accelerate AI adoption by eliminating operational bottlenecks. Pre-approved MCP servers and centralized policies enable faster deployment of AI agents and MCP servers without compromising security or compliance.



Safer AI Adoption

Empower your teams to leverage AI with confidence. Granular access controls and real-time monitoring ensure that AI agents can only perform authorized actions while maintaining complete auditability of all operations.



Scale Confidently

As your AI deployments grow across the enterprise, Barndoor scales with you. Centralized governance, automated policy enforcement, and unified logging ensure consistent security and compliance across all AI agents and MCP servers, regardless of scale.



Barndoor.ai

Ready For 2026?

As organizations move from experimentation in 2025 to operationalized, enterprise-wide AI in 2026, the need for secure, governed, and auditable AI access becomes urgent.

Barndoor transforms AI agent-MCP server access from a fragmented, high-risk model into a secure and governed process. By combining granular controls, real-time validation, and unified visibility, it empowers enterprises to harness AI's full potential, safely, responsibly, and at scale.

Interested in learning more?

Visit barndoor.ai

